

1. Brute-Force Solution

We can check each pair of coordinates (x_1, y_1) and (x_2, y_2) , calculate their local maximum area by multiplying the distance in x , $|x_2 - x_1|$, by the height of the smaller y , $\min(y_1, y_2)$. If the current pair has a greater local maximum area than the best found thusfar, we update variables to hold their indices and update max area. Checking each pair is $O(n^2)$.

2. Core Greedy Heuristic

In a sentence, we first try to maximize distance in x , and then work our way “inward,” shrinking x slightly, but seeking greater y by rejecting smaller heights. More specifically, we have left and right pointers which start at the left and right extremities respectively, and we increment the left pointer if the left’s y value is smaller, otherwise we decrement the right pointer. Each increment or decrement we recalculate a local maximum area, and compare to the best thusfar, updating our Answer if it’s better. We stop when the pointers meet.

3. Pseudo-code

```
1  findMaxRectangle(coordinates):  
2      best_thusfar = Answer(maxArea=0, left-x, right-x, height)  
3  
4      left_ptr = left-most-x, right_ptr = right-most-x  
5  
6      while (left_ptr != right_ptr):  
7          area =  $\min(\text{left\_ptr's } y, \text{right\_ptr's } y) * (\text{right\_ptr's } x - \text{left\_ptr's } x)$ ;  
8          if (area > maxArea of best_thusfar):  
9              update best_thusfar to hold area, the ptrs' x's, and min of the y's  
10             if (left_ptr's y was smaller):  
11                 increment left_ptr to next x to the right;  
12             else:  
13                 decrement right_ptr to next x to the left;  
14             return best_thusfar
```

4. Proof

The key insight is that when we start with a maximized x distance, the only way for any other rectangle to have a greater area is if it increases the height (since each subsequent x distance will be smaller or equal to the original x -maxed rectangle). And the only way to increase the height of a rectangle is to increase the smaller of the two y values. Therefore, when we start with the maximized

x , the only way any other rectangle can have a larger area is if the smaller y changes. This is true on each subsequent iteration as well – the only way to possibly achieve a greater area than the previous iteration (since each iteration shrinks the x distance) will be for the smaller y to change, so we move the corresponding pointer to avail ourselves to that possibility. To sum up, since we start with a given area, and continuously take steps that represent the only possible way to find a larger area than the previous step, we will find the area that is maximized as much as possible, thus solving the problem.

Performance-wise, it's easy to see that this is $O(n)$, because each iteration reduces the distance in x between the pointers by 1, and this distance goes from $n - 1$ at the beginning to 0 at termination (when the pointers meet).